

6. Optic flow estimation

Abstract

Optic flow estimation is an important building block for applications where visual information change dynamically over time. Mobile robot navigation, traffic monitoring, and human action recognition are three examples, among many others, which require or can benefit from optic flow information. In this lab, we will implement one simple, but foundational method for optic flow computation, the classical Lucas-Kanade algorithm.

Keywords

Optic flow • Lucas-Kanade algorithm

Contents

| | | |
|-----|------------------------------------|---|
| 1 | Lucas-Kanade algorithm | 1 |
| 1.1 | Computing the gradients | 1 |
| 1.2 | Computing $A^T A$ and $A^T b$ | 1 |
| 1.3 | Solving for the flow | 1 |
| 2 | Synthetic motion | 1 |
| 3 | Real sequences | 2 |
| 4 | Additional and optional activities | 2 |

1. Lucas-Kanade algorithm

Our first task is to complete the function

```
u,v = optical_flow(I1, I2, window_size, tau)
```

to compute the optic flow (u, v) between the image pair (I_1, I_2) with Lucas-Kanade (LK) algorithm, using local square windows of side length w (`window_size`), and τ (`tau`) as the minimum value of the eigenvalues of the matrix $A^T A$. Let's proceed step by step.

Computing the gradients

We first smooth the input images with a Gaussian filter, and then we have to compute the spatio-temporal gradients, I_x , I_y and I_t on the smoothed images. To that end, we define the corresponding three kernels for the x , y and t directions, and then apply 2D convolutions. The kernel and convolution for x is given. **1 ► Now, write the code to define the two other kernels and apply their respective convolutions.** For the temporal gradient, follow this: $I_t = (I_2 - I_1) * M_t = I_2 * M_t - I_1 * M_t$, where M_t is the 4×4 mask for a mean filter. Notation reminder: we use $*$ to denote the convolution operation (not a matrix product or an elementwise product), and I_h represents the image derivative with respect to the axis $h \in \{x, y, t\}$.

Computing $A^T A$ and $A^T b$

Remember that in LK method, for each local window of n pixels, we have that $Au = -b$, where A is the $n \times 2$ matrix of

the n spatial gradients, and b is the $n \times 1$ vector of temporal derivatives. To find out the optic flow u for that window we have that $u = (A^T A)^{-1} A^T b$.

In the provided code, you are given the arrays I_x , I_y , and I_t that contain the corresponding derivatives just for the local window around a particular pixel at (i, j) . From these, **2 ► you only have to compute the variables $A^T A$ (for $A^T A$) and $A^T b$ (for $A^T b$).**

Solving for the flow

Finally, **3 ► we can use `linalg.lstsq(C,d)`** to compute the least-squares solution x to $Cx = d$. Just pay attention to what matrix C , vector d , and x are in our case.

Remember that we can only compute the optic flow in those image locations where $A^T A$ is well-conditioned. So, before solving for u we must check for this. **4 ► We can verify** that either the smallest eigenvalue for $A^T A$ is larger than τ , or that $A^T A$ has rank 2. For this, we can rely on `np.linalg.eigvals()` and `np.linalg.matrix_rank()`, respectively.

2. Synthetic motion

We now have to test how `optical_flow()` works. Since the basic LK algorithm assumes small motions, it is good that we can generate the motions we want, under controlled conditions. It is harder to find real sequences that have the motion we are interested to check.

As you can see in the code, we use `SimilarityTransform()` to define the transformation with the desired amounts of translation, scale and rotation, which is then applied to the actual image with `warp()`.

To visualize the computed flow, you are provided with the function `display_optic_flow()`. Good visualisations of outputs of algorithms of optic flow (and others) is always critically important during the development and analysis of these algorithms.

In the following, you can use the threshold $\tau = 0.01$. **5 ► Experiment with different amounts of translation** and find out up to which magnitudes LK fares well. **6 ► Regarding**

what the function `plot_optic_flow()` displays, make sure you understand, besides the arrows, what the colors in the background represents. Check the results look qualitatively good (when translation is small enough). 7 ► *Complement this visual inspection* of the flow map with the histogram of the magnitude of the optical flow. 8 ► *Vary also the window size w* (from the smallest $w = 3$ up to a few tens) and notice the impact on the estimated flow (both the map and the magnitude histogram). Does w have an influence on the sparsity/density or smoothness of the map? Try to find an explanation.

9 ► *Repeat the above experimentation* with the scale factor (both “zoom in” and “zoom out”), and then with rotations. Notice that, unlike the translations, the magnitude of the flow at different image positions should be different. Make sure the output make sense, qualitatively.

3. Real sequences

We can now study how optic flow performs in some image sequences from real scenes. Although ground-truth optic flow is hardly available in these cases and, therefore, estimation error cannot be easily computed, we can anyway evaluate the result qualitatively in these more challenging (and interesting) scenarios, where one or more objects may move independently, in contrast to the global movement of the previous synthetic examples.

10 ► *Experiment on some of consecutive frames* of the Hamburg taxi sequence, which is one example of the classical sequences which optic flow methods used to be tested on. Although local motion in this sequence can be considered to be “large” and, therefore, the vanilla LK method might not perform very well, find for which window sizes w we can get reasonable results. What potential benefit can we get by increasing τ one or two orders of magnitude?

4. Additional and optional activities

Remember that all exercises are optional, and those marked explicitly as “optional” are particularly so. We propose many optional exercises because we do not want you to be left wanting more activities. Of course, this does not mean we want you to try them all, let alone complete them all successfully. Therefore, be wise and selective in which exercises you want to try and how much time/effort you are willing to invest, depending on your background, your available time, your goals...

11 ► *(Optional) Compute the Gaussian pyramid of the input images* (you can use `pyramid_gaussian`). Experiment a little bit on the estimation of flow at different levels of the pyramid: as you know, the optic flow becomes smaller at downscaled versions of the images, which is the basic observation for the coarse-to-fine version of the LK algorithm.

12 ► *(Optional) Try to approximately segment the object(s) in real sequences* using optical flow information. The idea is to use one simple segmentation algorithm that you know (maybe Otsu), but instead of applying it on gray-level values, we can use optic flow. After the segmentation, we can find the regions (connected components). For instance, can we isolate the different vehicles in the taxi sequence?

13 ► *(Optional) Write and test a function to display* several optic-flow information (e.g., u , v , orientation and magnitude). You may take this opportunity to learn about the flexibility of arranging multiple Matplotlib axes in a grid that is offered by `GridSpec`.

14 ► *(Optional) Study* the `quiver()` function to display the output of optic flow, which is used in our `display_optic_flow()`.

15 ► *(Optional) Learn about about a commonly used color wheel representation* which combines the magnitude and orientation into a single 2D map. You can find examples in, for instance, [this paper](#). Try some available implementation of this color coding, or implement it yourself.

The following exercises are more open, and possibly harder and significantly more time consuming. Some are even advanced. Therefore, you should only consider them if you find them, or the topic of this lab, motivating or interesting, and you have lot of spare time to spend on them. After all, they are less relevant in the context of our introductory course.

16 ► *(Optional, advanced) Implement and experiment with a hierarchical, coarse-to-fine version* of LK. We can expect that larger motions can be estimated.

17 ► *(Optional) Try the Python implementation* of the hierarchical LK in OpenCV on some sequences (the ones used in this lab and/or others) and compare with our basic (single-level) implementation. You may also like to follow [this tutorial](#).

18 ► *(Optional) Consider one of the full real sequences*, and generate either an on-screen animation (with Matplotlib), or a nice-looking video of some representation of the computed flow, either alone or overlapped with the frames.

19 ► *(Optional, advanced) Study some other well-known and good optic flow method*. For instance, for the mathematically inclined, in the pre-deep-learning era, [variational methods](#) were popular in computer vision. [Brox et al. \(ECCV 2004\)](#) proposed a variational algorithm for optical flow. If you wish, find also some available implementation of this (or other) approach, and experiment a little bit with it. You should at least observe that this better methods outperform our simple LK implementation.

20 ► *(Optional, advanced) Find some (small) dataset* with available optic flow ground-truth. Find out which performance metrics are commonly used to assess the performance. Then, you can run LK under different w , or other OF methods, and compare the results against the ground-truth using some of those metrics.